



# Kubernetes for Data Engineers

Big Data on Kubernetes - [Day 1]



LUAN MORENO  
CEO & Data Architect  
Data Engineer & MVP

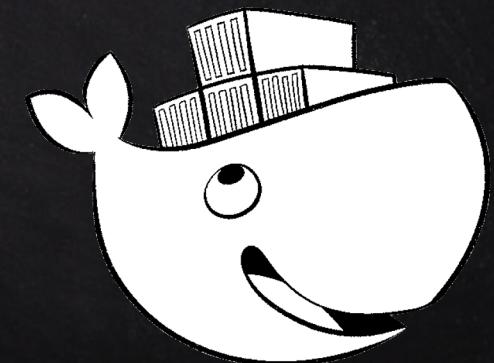


MATEUS OLIVEIRA  
Big Data Architect  
Data In-Motion Specialist



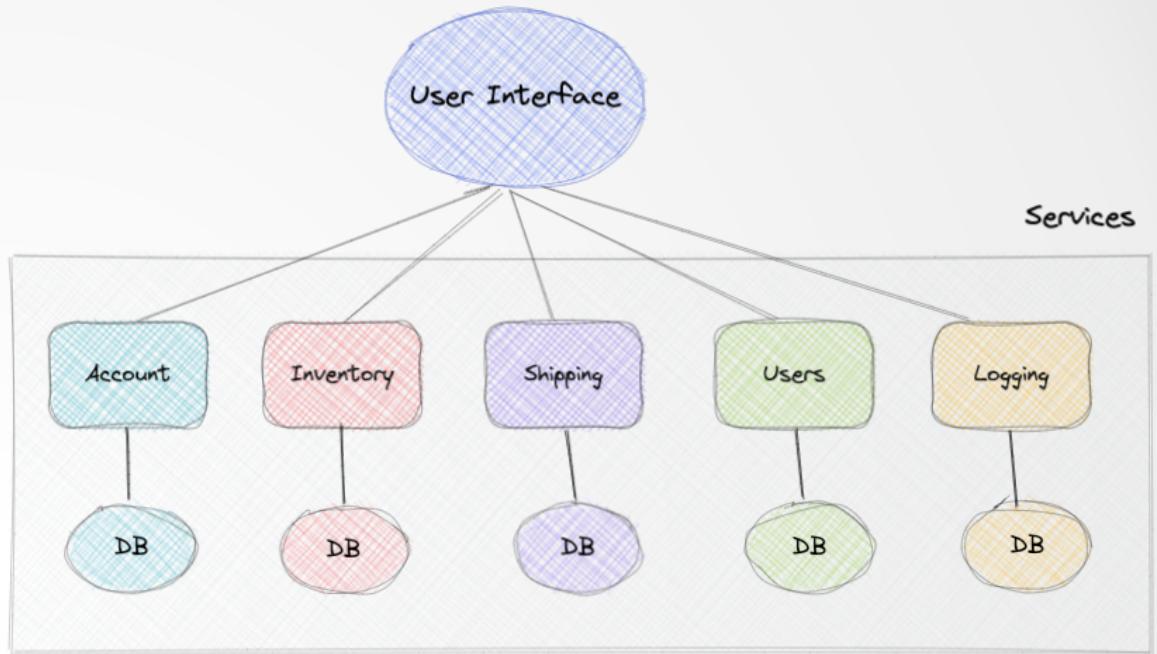
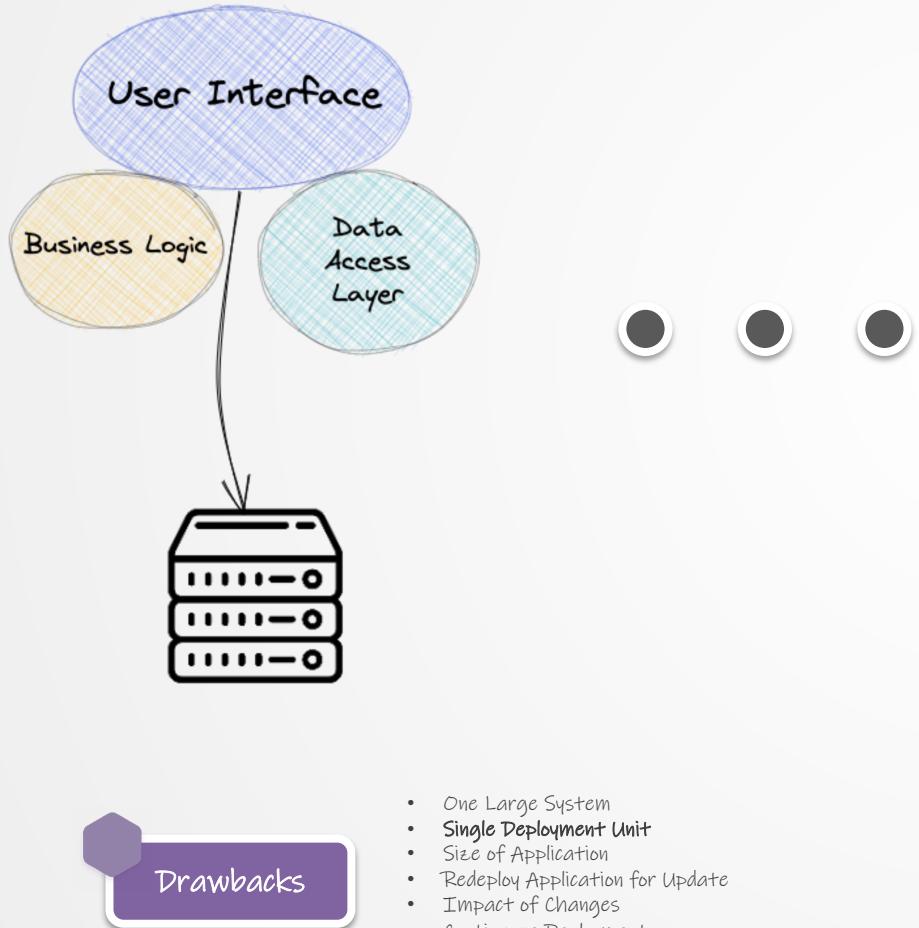
# *History of Containerized Applications and Container Orchestration Tools*

- Monolithic vs. Microservices
- The Era of Microservices & Big Data
- The Principles of Containerized Applications
- Stateless vs. Stateful Applications
- Container Orchestration Engines



# Monolithic vs. Microservices

Single Unit Deployment vs. Smaller Independent Units for Speed Up Development Cycle



- Decomposing Application
- Manageable Services for Faster Development
- Develop Independently Behavior
- Adopting New Technologies
- Scale Components Individually

- Distributed System
- Different Architecture Paradigms
- Performing Transactions
- Testing Application
- Deployment of Service

# The Era of Microservices & Big Data

Real-Time Analytics at Scale by Enabling Services to Communicate in a Unified and Central Hub

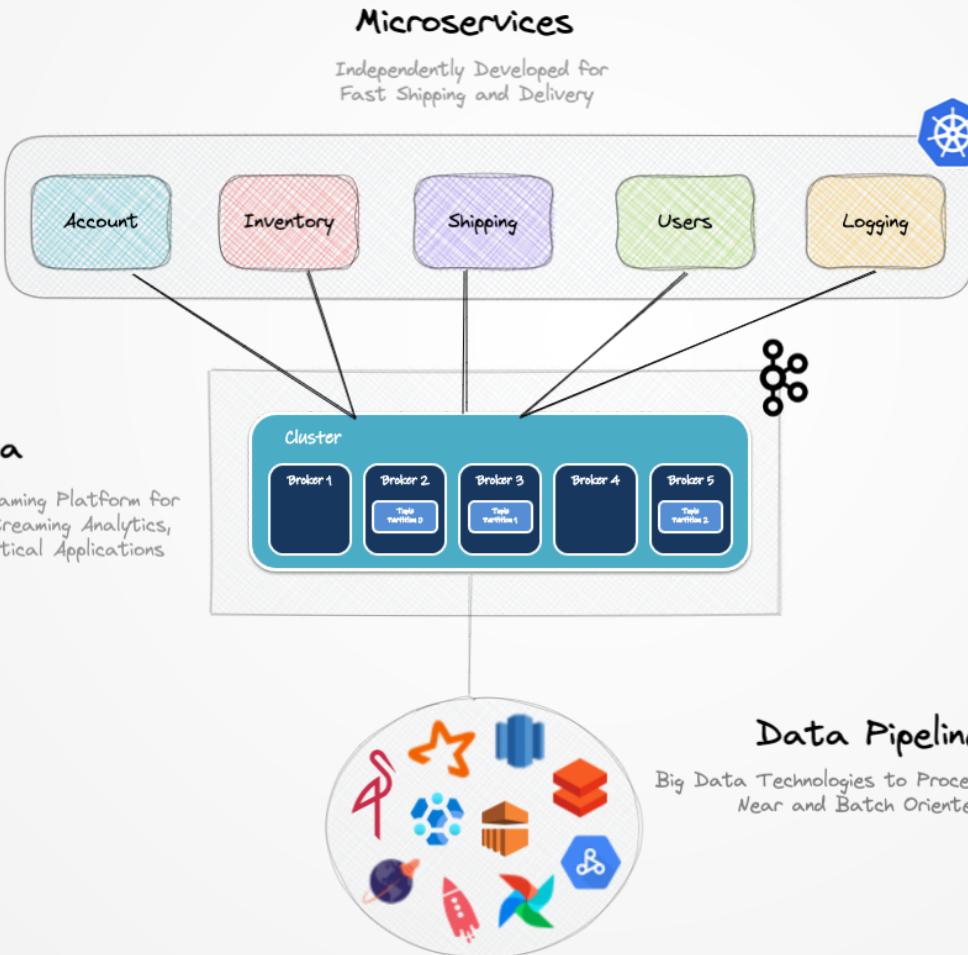


## Service Mesh

- Software Architecture
- Dedicated Infrastructure Layer
- Service-to-Service Communication
- Sidecar Proxy
- Service Discovery
- Load Balancing
- Encryption
- Authentication & Authorization
- Monitoring and Logging
- Circuit Breaker Pattern

## Apache Kafka

Open-Source Distributed Event Streaming Platform for High-Performance Data Pipelines, Streaming Analytics, Data Integration, and Mission-Critical Applications

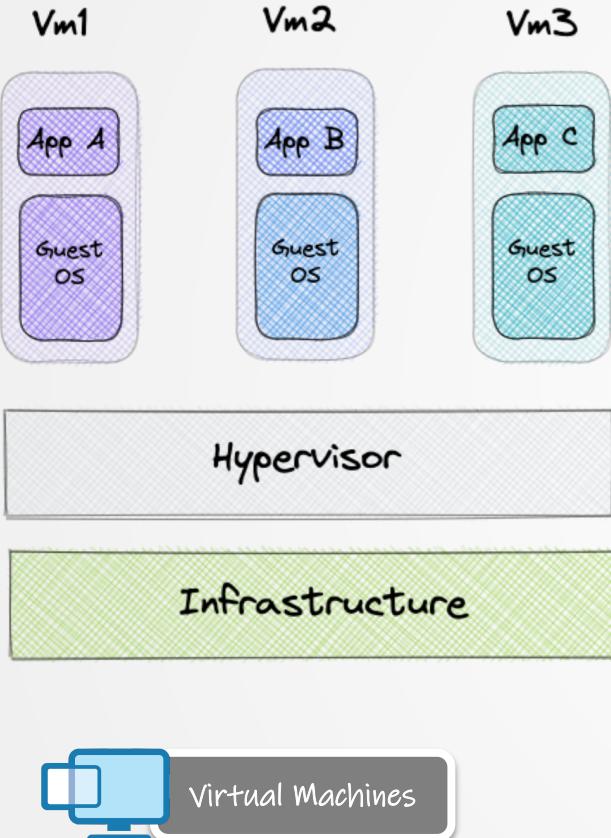


## Data Mesh

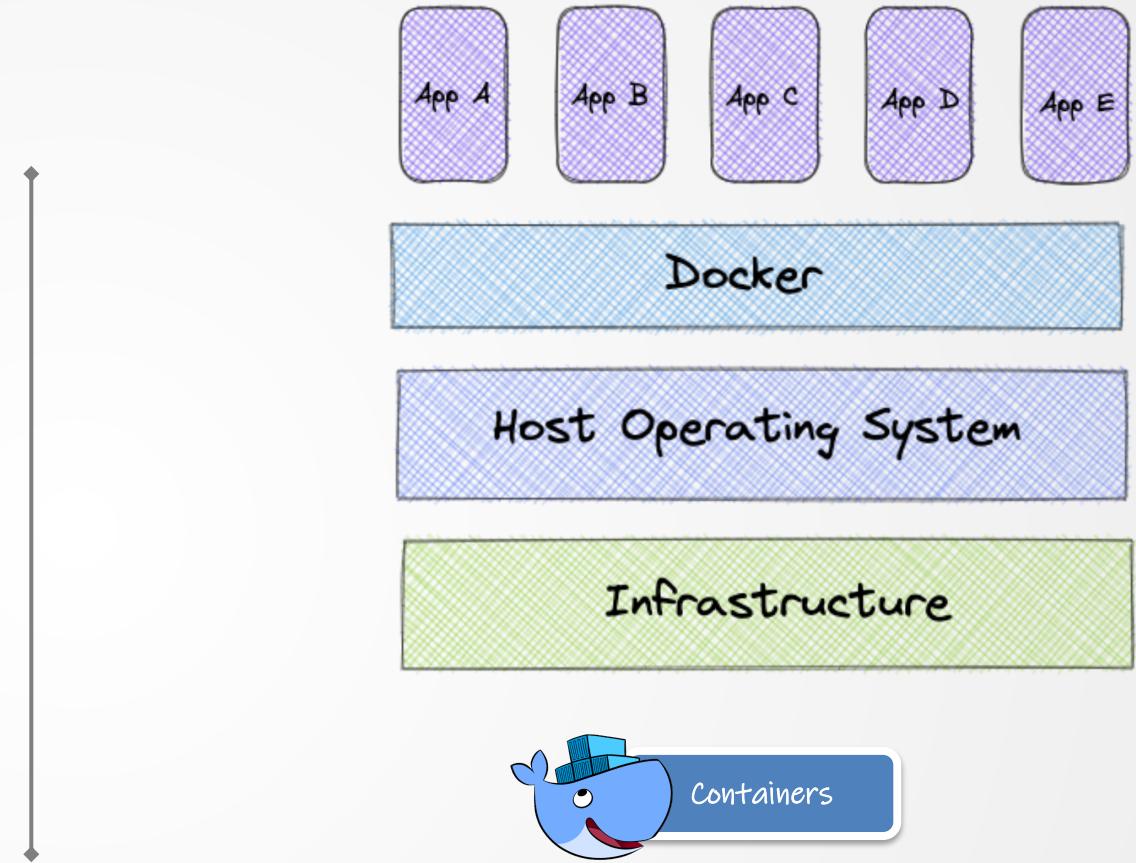
- Design Modern Data Architectures
- Data-Centric and Data Management
- Data Ownership by Domain (1)
- Data as a Product (2)
- Data Everywhere Self-Serve (3)
- Data Governed (4)
- Specific Business Domain

# Virtual Machines vs. Containers

VM - Abstracts Machine and Containers - Abstracts Applications



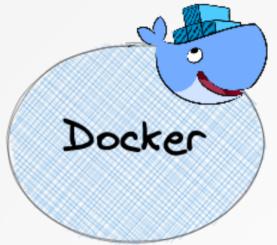
- Enables Virtualization Process = Hypervisor
- Allow Multiple OS
- Ineffective Resource Usage
- Dependent of a Guest OS
- Not Portable



- Images from MB to GB
- Portability
- Effective Resource Usage
- Easier for Maintenance
- Highly-Scalable
- The State of Art

# Docker

Use OS-Level Virtualization to  
Deliver Software in Packages



13 Million +  
Developers

7 Million +  
Applications

13 Billion +  
Monthly Image Downloads

## Docker Desktop

Developer Productivity Tools and a  
Local Kubernetes Environment

## Docker Hub

Cloud-Based Application  
Registry and Development  
Team Collaboration Services



## Tools

- Docker Compose = Running Multi-Container Docker Applications
- Docker Swarm = Native Clustering Functionality
- Docker Volume = Copy or Create a File in a Container



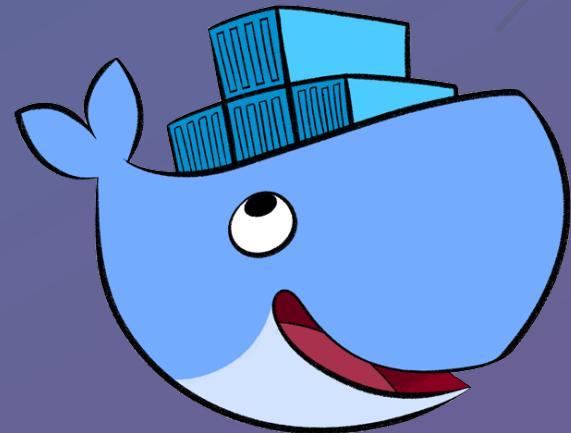
## Information

- First Started in 2013, Developed by Docker Inc
- Platform as a Service (PaaS)
- OS-Level Virtualization



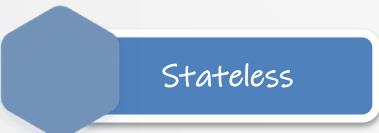
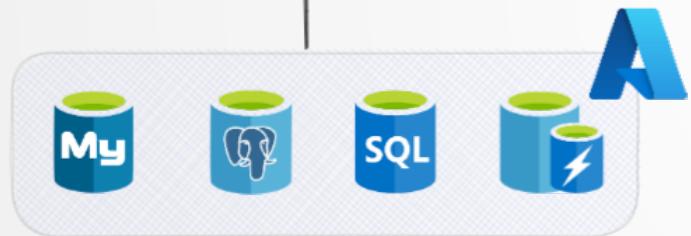
# Docker 101

*Demo*



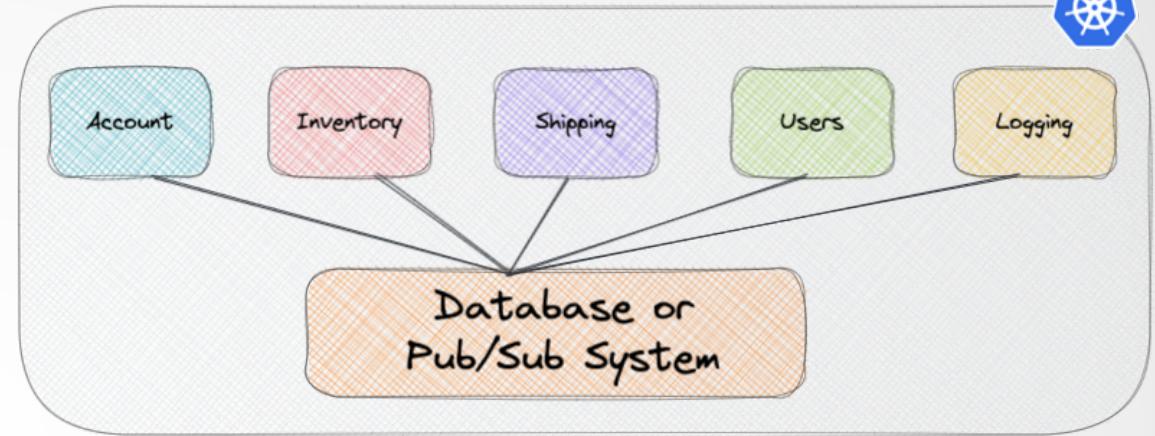
# Stateless vs. Stateful Applications

Stateless - Does Not Save Client Data Generated and Stateful  
- Save Data and State



- Design, Implementation and Architecture is Simple
- Handles Crash Efficiently
- Scaling is Fast and Easy
- **Recommendation** - Kubernetes and Istio

Store Data in Persistent Disk Storage Systems  
on Kubernetes and Store User Data



- Complex Design and Architecture
- Storage System Design is Key
- Resource Isolation for Better SLA
- Used by Modern Data Apps
- **Recommendation** - Kubernetes and YugaByteDB



# Container Orchestration Engines

Automates Deployment, Management, Scaling and Networking of Container Engines



Docker Swarm

- Maintained by Docker
- Multiple Docker Instances
- Tied with Docker API
- Used to Spin Up Multiple Containers



OpenShift

- Developed by Red Hat
- Hybrid Proposal
- Commercial Product
- Kubernetes Encapsulated
- Uses CRI-O Runtime
- Image Repository Out-of-Box



Kubernetes

- First Open-Source Orchestration Engine
- Originally Designed at Google
- Maintained by CNCF Corporation
- Runtimes - Docker, Containerd, CRI-O
- Container Runtime Interface (CRI) - Standard

# *Kubernetes Concepts*

- Introduction and Overview
- Cluster Architecture
- Master & Worker
- Workloads
- Service
- Ingress
- Storage
- Configuration



# Introduction & Overview

Most Used Orchestration Engine Nowadays. Deployed at Google's Infrastructure, Azure and AWS Data Centers

De-Facto Operating System  
for Cloud Native Products



## Characteristics

- Originally Developed by Google (Borg & Omega)
- First Open-Source Orchestrator Platform - 2017
- Self-Healing Infrastructure Backbone System
- Velocity, Scaling, Abstraction and Efficiency



IaaS, PaaS and SaaS Softwares  
for Easy and Fast Development

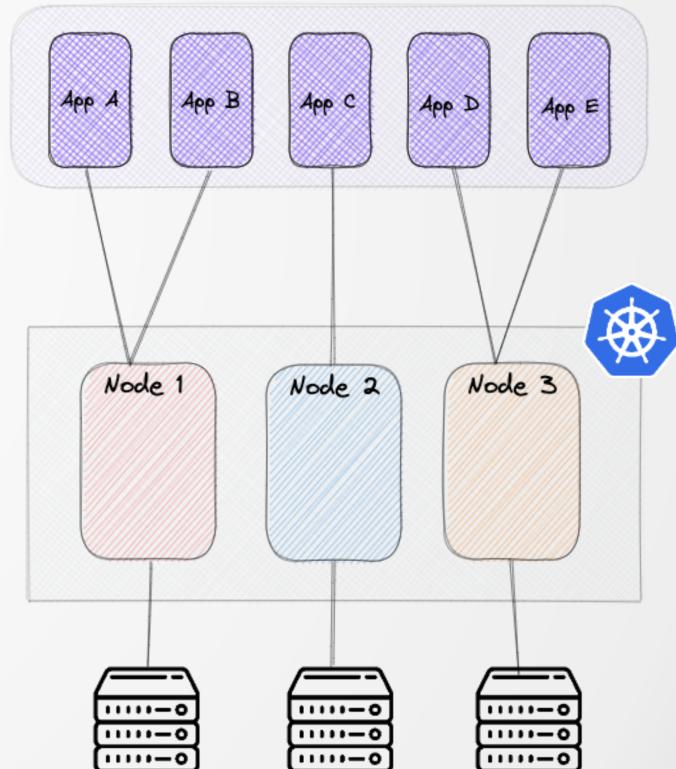
## Main Features

- Service Discovery & Load Balancing
- Storage Orchestration
- Automated Rollouts and Rollbacks
- Automatic Bin Packing
- Secret & Configuration Management

Revolution  
The Cloud Native World



- \* First - Creation of Cloud Computing
- \* Second - Dawn of DevOps
- \* Third - Containers

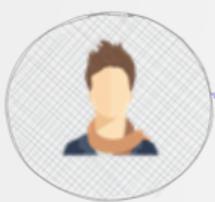


# Cluster Architecture - Control Plane and Nodes

Consists of a Set of Worker Machines (Slave) Called Nodes that Runs the Containerized Applications and an Intelligent Control Plane Area

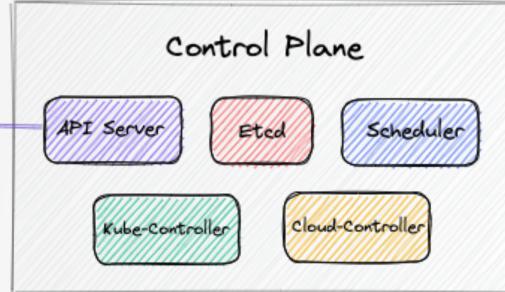
## Users

- \* API
- \* CLI
- \* UI



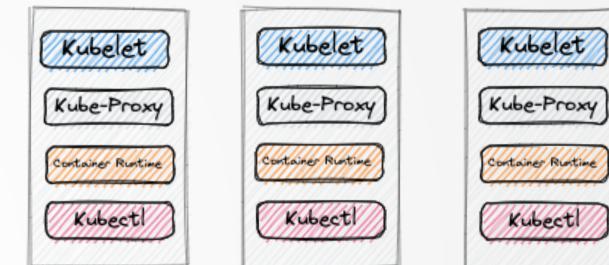
## Master Node

Controls and Coordinates Cluster Activities



## Worker Nodes

Runs Pods and Applications



## Factors Taken

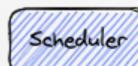
- Individual and Collective Resource
- Hardware, Software
- Policy Constraints
- Affinity and Anti-Affinity
- Data Locality
- Inter-Workload Interference
- Deadlines



- Frontend Server for Handling API Requests



- Key-Value Store Database



- Watcher for Newly Created Pods with no Assigned Node



- Running Resource Controller Processes - Deployments



- Interacts with Cloud Providers - LBs, Volumes



- Agent Running on Each Node. Checks Container Health and State



- Networking to Route Request Between Pods on Different Nodes and Between Pods & Internet



- Starts and Stops Containers and Handles Communications, Usually Docker (Rkt, CRI-O)



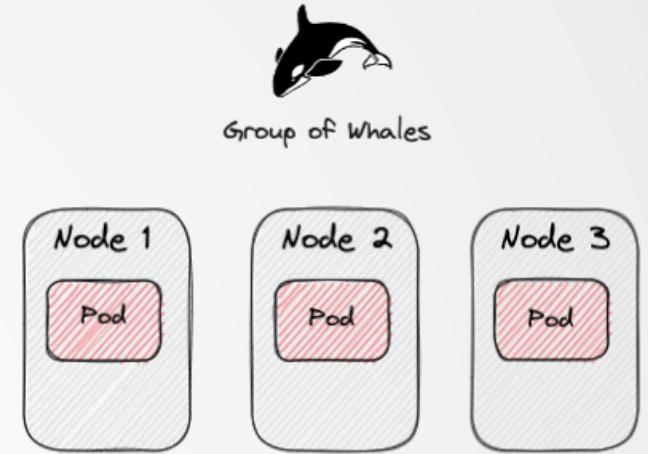
- Command-Line Tool for Interacting with Kubernetes Objects

# Pod & Manifests

Smallest Execution Unit in Kubernetes, Ephemeral by Nature, Automatically Create a New Replica Transparently to Continue Operations



```
apiVersion: v1
kind: Pod
metadata:
  name: mc1
spec:
  volumes:
    - name: html
      emptyDir: {}
  containers:
    - name: 1st
      image: nginx
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
    - name: 2nd
      image: debian
      volumeMounts:
        - name: html
          mountPath: /html
      command:
        - /bin/sh
        - '-c'
      args:
        - while true; do date >> /html/index.html; sleep 1; done
```



Phases

- **Pending** = Accepted by Cluster, Waiting, Downloading or Waiting Process
- **Running** = Bound to a Node, All Containers Created Successfully
- **Succeeded** = All Containers in Pod have Terminated in Success
- **Failed** = Exited with Non-Zero Status or Terminated by the System
- **Unknown** = State was not Obtained, Occurs due to Error in Communicating with Node

# Workloads

Applications Running, Singles or Several Components Working Together Inside of Pods

Specify How Many Identical Pods are Needed

## Deployment & ReplicaSet

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
```

Groups of Replicated Pods (Copy in All Nodes)

## DaemonSet

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd-elasticsearch
spec:
  selector:
    matchLabels:
      app: fluentd-elasticsearch
  template:
    spec:
      containers:
        - name: fluentd-elasticsearch
...
...
```

Pods with Sticky Identity and Persistent Storage

## StatefulSet

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  selector:
    matchLabels:
      app: nginx
  serviceName: "nginx"
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
  spec:
    terminationGracePeriodSeconds: 10
    containers:
      - name: nginx
        image: k8s.gcr.io/nginx-slim:0.8
        ports:
          - containerPort: 80
            name: web
    volumeMounts:
      ...
    volumeClaimTemplates:
      - metadata:
          name: www
        spec:
          accessModes: [ "ReadWriteOnce" ]
          storageClassName: "my-storage-class"
          resources:
            requests:
              storage: 1Gi
```

Runs a Pod for a Specified Number of Times

## Job

```
apiVersion: batch/v1
kind: Job
metadata:
  name: queue-worker
spec:
  completions: 1
  parallelism: 10
  template:
    metadata:
      name: queue-worker
    spec:
      containers:
```

Schedule Jobs by Cron Daemon

## Cronjob

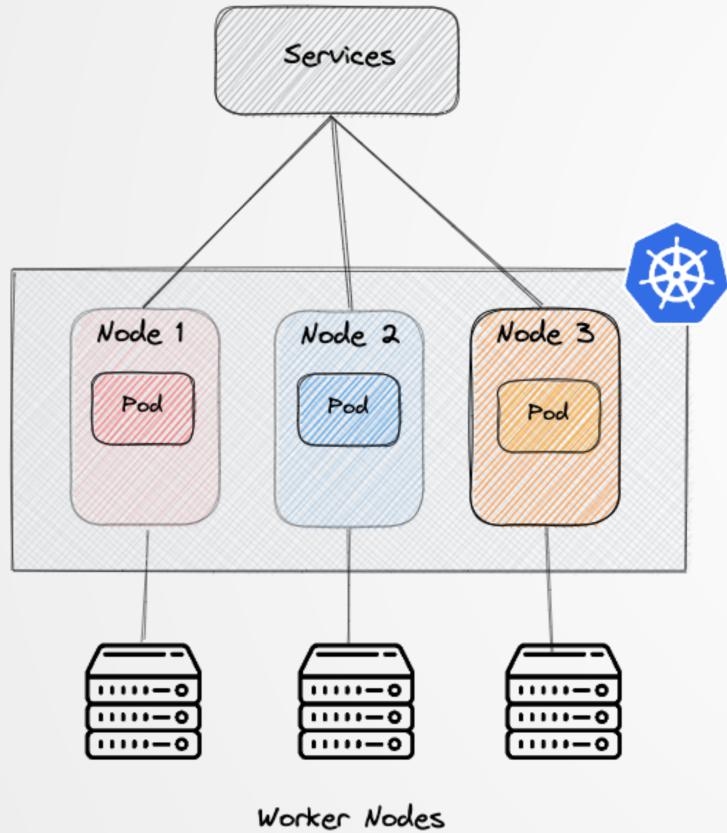
```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: demo-cron
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec: ...
```

## Pod Controllers

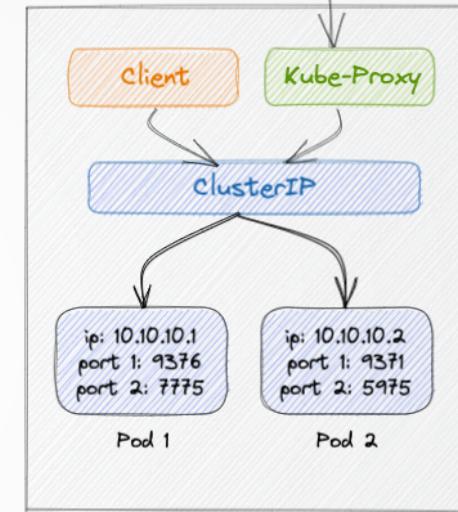
Built-In Workload Resources

# Service

Abstract Exposure of an Application Running on a Set of Pods as a Network Service. Pods Receives IP Addresses and a Single DNS Name



Control Plane



Worker Node



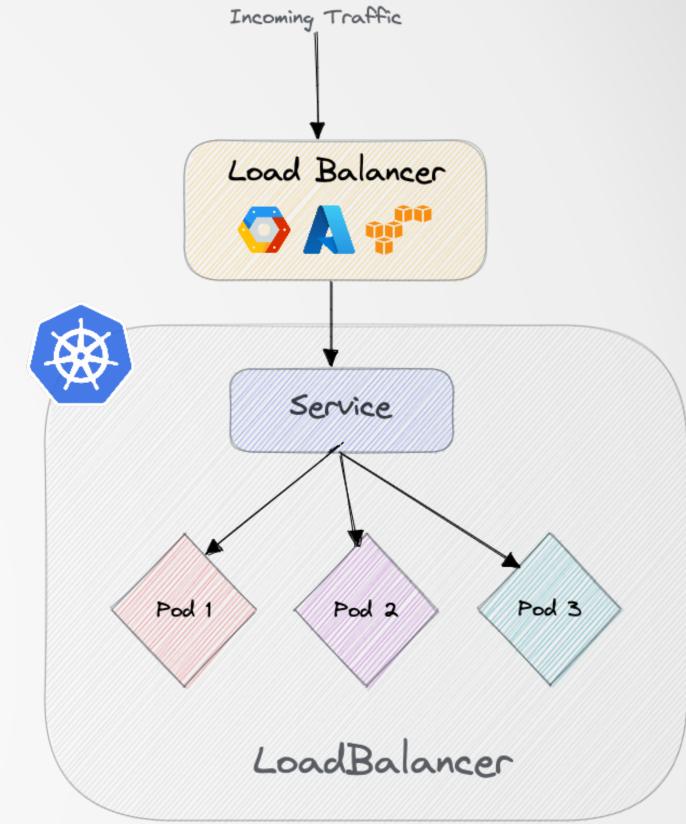
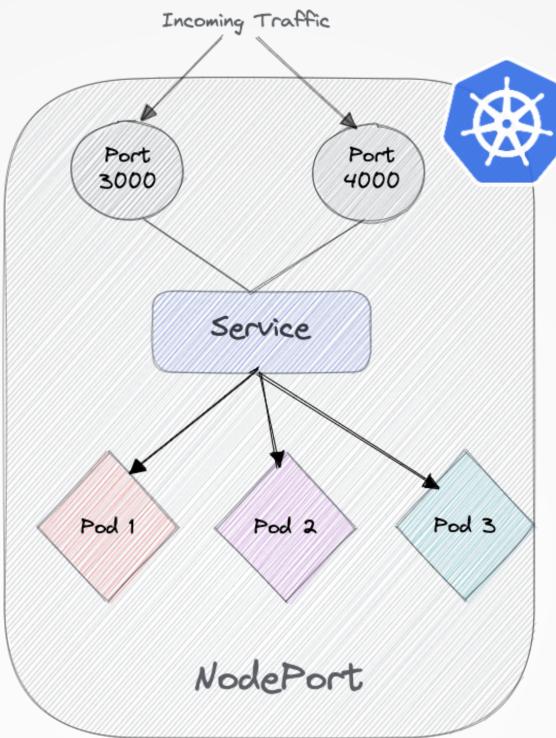
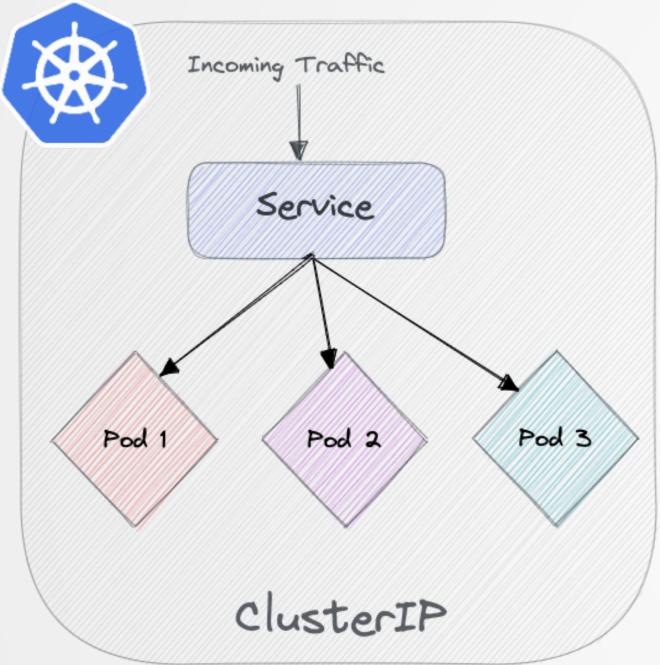
Selectors = Identify a Set of Objects

- Equality and Set-Based
- Multiple Requirements

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

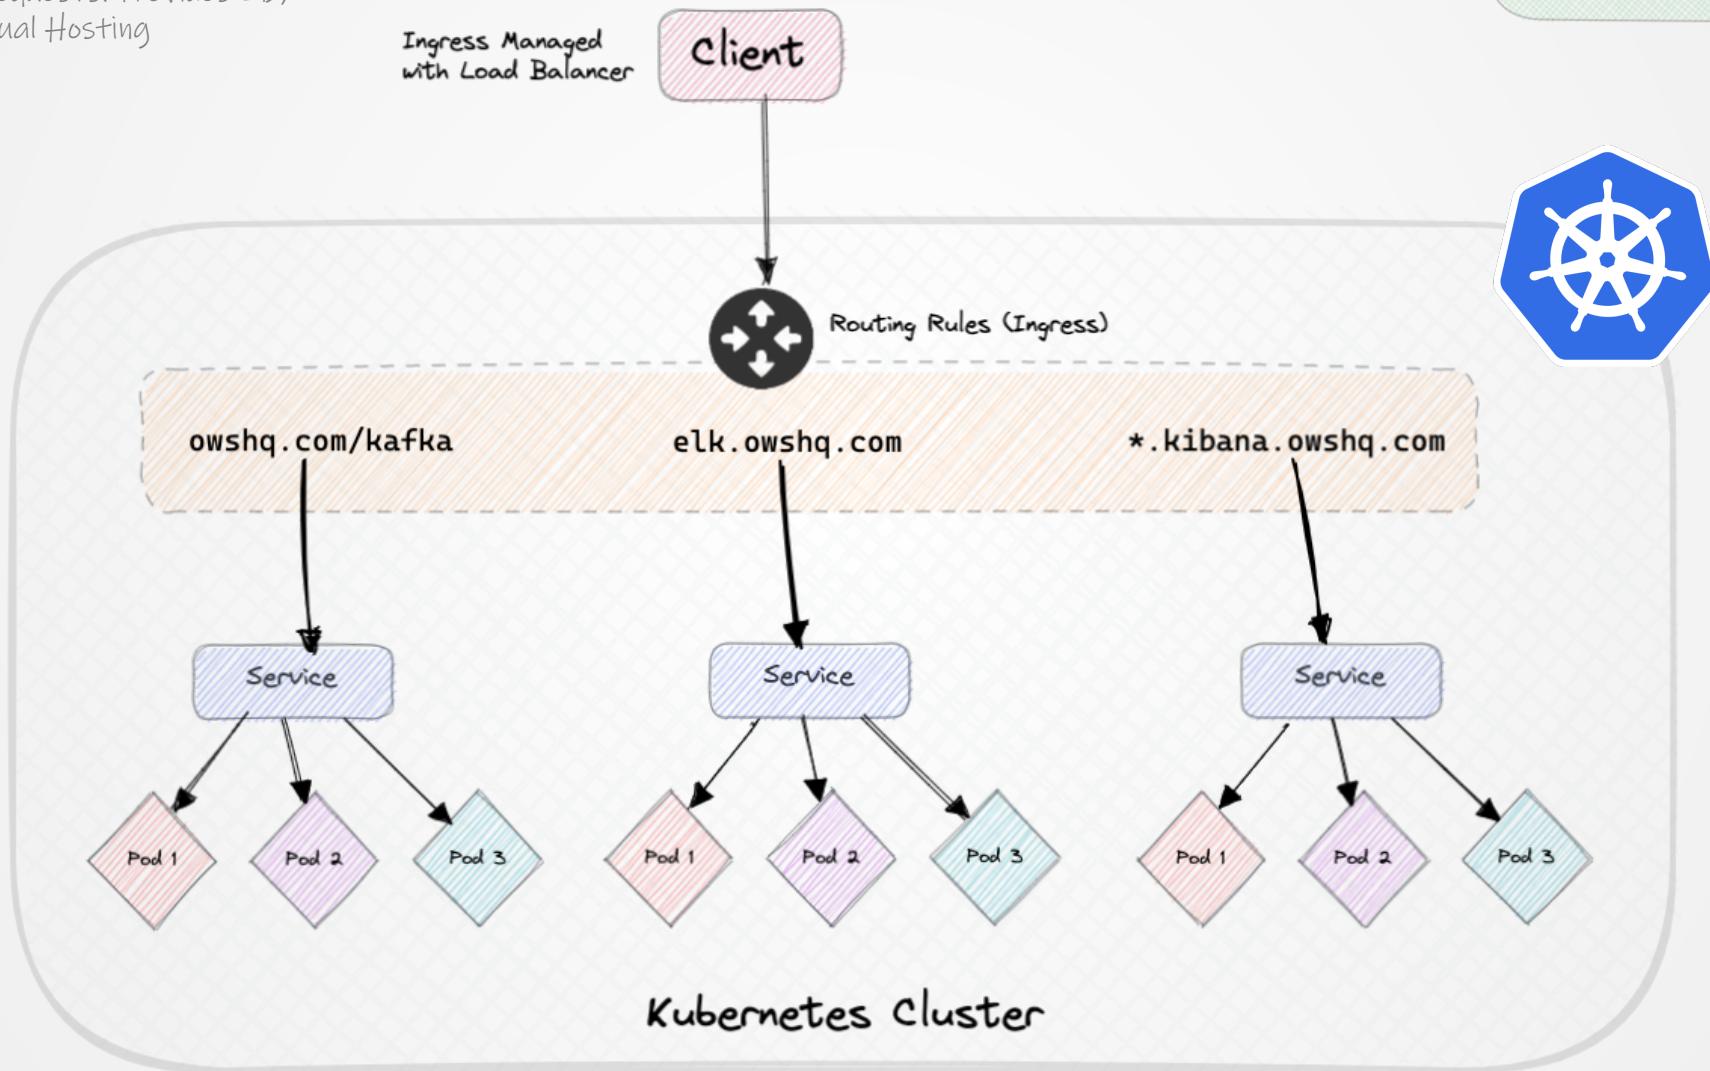
# Service Types

Expose a Service onto an External API Address, Outside of Kubernetes Cluster for Internet-Based Access



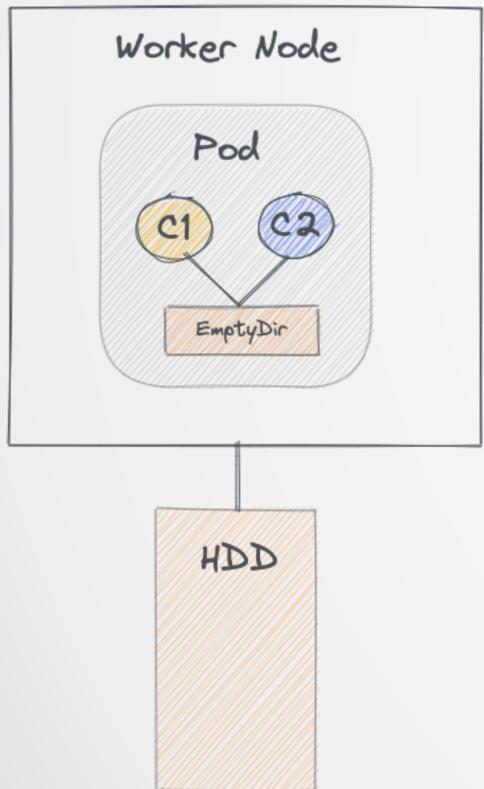
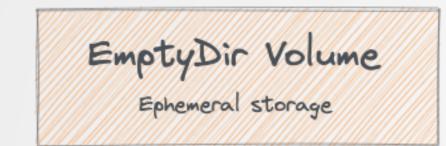
# Ingress

Manages External Access to Services in a Cluster, Typically HTTP Requests. Provides LB, SSL and Name-Based Virtual Hosting



# Volumes & Types

Pod Restarts in Clean State. Wipe Out Data. Ability to Use Different Volumes Simultaneously. Destroys Ephemeral but do not Destroy Persistent Volumes



## Ephemeral Types

Data not Stored Persistently Across Restarts

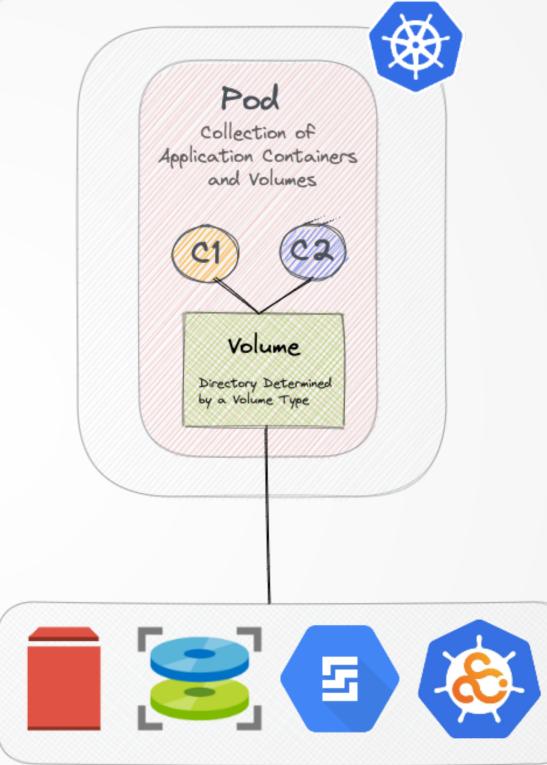
- \* emptyDir
- \* CSI ephemeral volumes
- \* generic ephemeral volumes

```
apiVersion: v1 kind: Pod  
...  
spec:  
  volumes:  
    - name: cache-volume  
      emptyDir: {}  
  containers:  
    - name: demo  
      image: cloudnativized/demo:hello  
      volumeMounts:  
        - mountPath: /cache  
          name: cache-volume
```



## Worker Node

Physical or Virtual Machines



## Volume Type

Directory Determined by a Volume Type

- \* AWS EBS
- \* Azure Disk
- \* Google Persistent Disk
- \* Container Storage Interface

## Container Storage Interface (CSI)

Standard for Exposing Arbitrary Block and File Storage Systems for Container Orchestration Systems (COs)

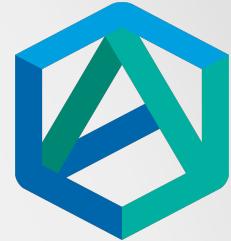


CSI = Container Storage Interface

- Promoted GA in v1.13 Release
- Powerful Volume Plugin System
- Before Embedded on Binaries
- Expose Arbitrary Block and File Storage
- Enable Dynamic Provisioning



# Configuration - ConfigMaps & Secrets

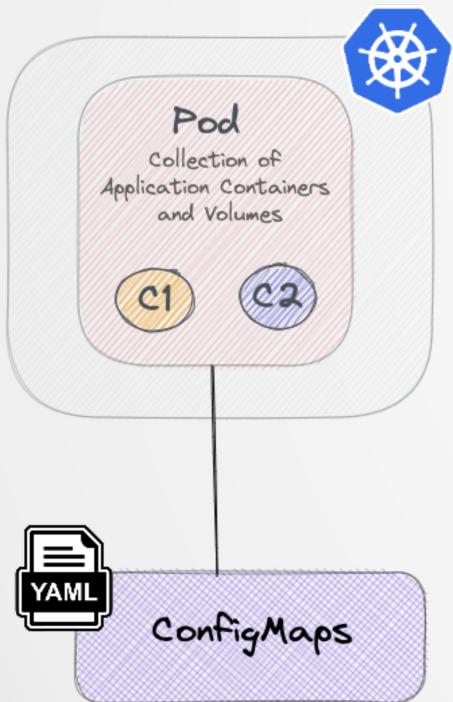


ConfigMaps = Used to Store Non-Confidential Data in Key-Value

Secrets = Sensitive Information

## ConfigMaps

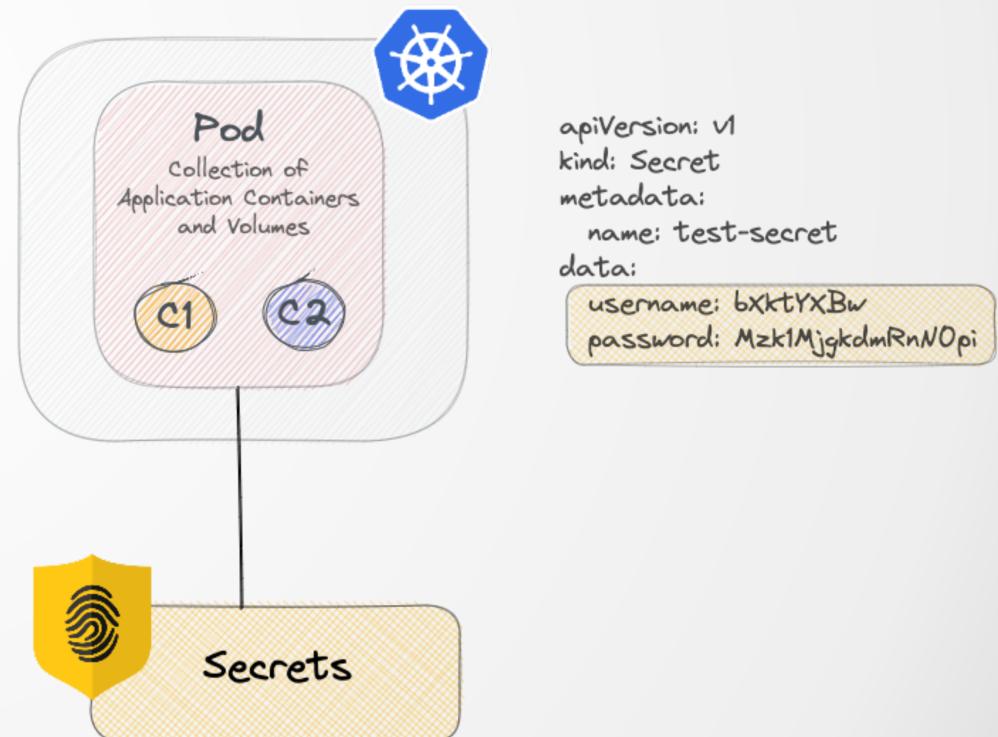
Store Non-Confidential Data in Key-Value Pairs  
Consumed as Variables, CLI Arguments or  
as Configuration Files in a Volume



```
kind: ConfigMap
metadata:
  name: game-demo
data:
  player_initial_lives: "3"
  ui_properties_file_name: "user-interface.properties"
  game.properties: | 
    enemy.types=aliens,monsters
    player.maximum-lives=5
  user-interface.properties: |
    color.good=purple
    color.bad=yellow
    allow.textmode=true
```

## Secrets

Sensitive Data such as Password, Token or a Key.  
Store Confidential Data Outside of Application Code



```
apiVersion: v1
kind: Secret
metadata:
  name: test-secret
data:
  username: bxktyXBw
  password: Mzk1MjgkdmRnNOpi
```

# Kubernetes 101

*Demo*



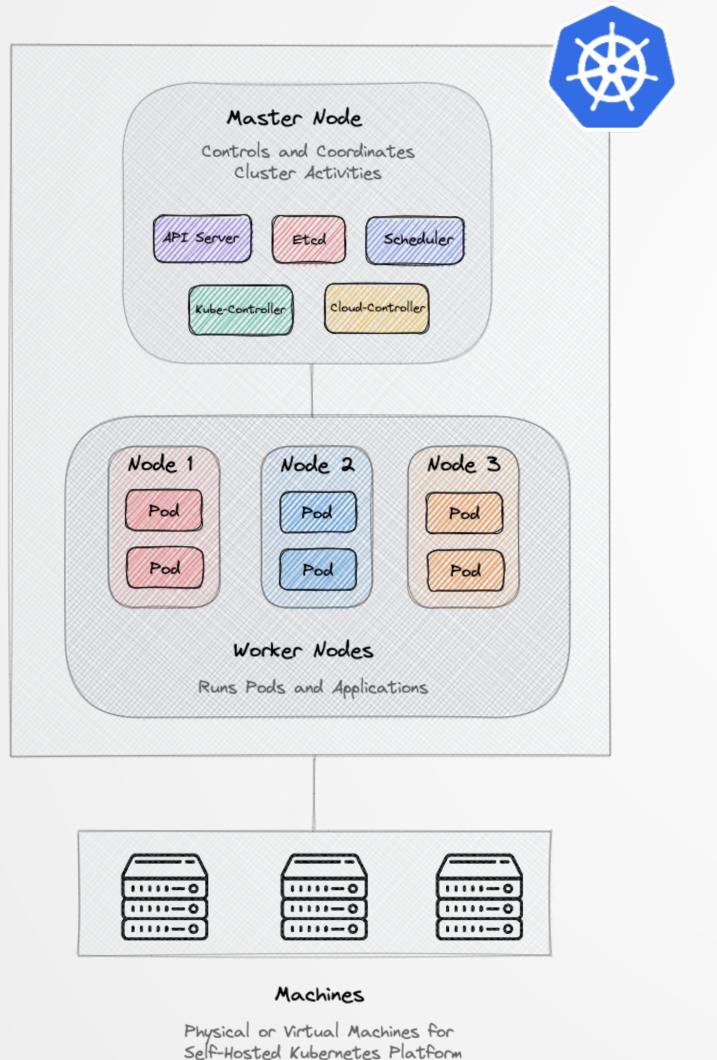
# *Kubernetes as a Managed Offering*

- Self-Hosted Kubernetes
- Managed Kubernetes Services
- AKS, EKS and GKE
- Infrastructure-as-Code (IaC)



# Self-Hosted Kubernetes

Usually Deployed in On-Premises Environments,  
Hard to Maintain and Implement



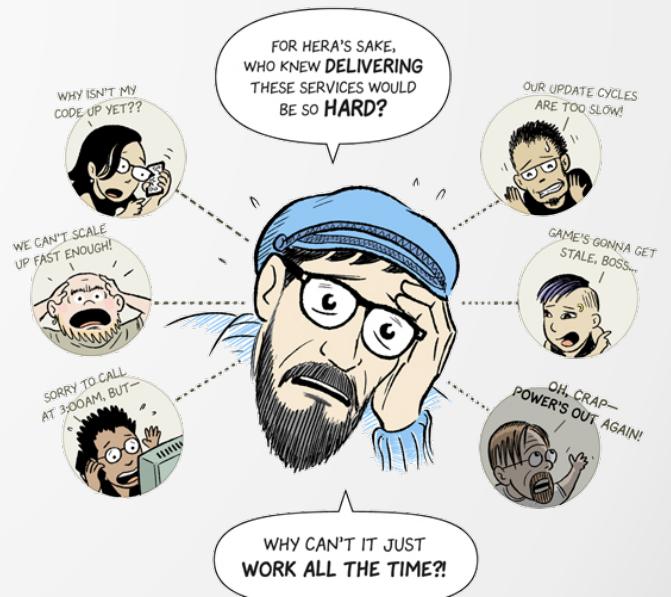
## Challenges

- High Availability
- Single Sign-On
- Multitenancy
- Resource Isolation
- Permission Management
- Upgrades
- Backups
- Package Management
- CI/CD Integration



## Implement Kubernetes Self-Hosted from Scratch

- \$1 Million Dollars in Engineer Salary
- Still Undergoing Process
- 6 Team Members Upstream Production
- 6 Months of Work



# Kubernetes as a Managed Offering

Managed Production-Ready Solution Provided by Cloud Providers  
for Deploying Your Containerized Applications

certified



kubernetes

## Azure Kubernetes Service (AKS)

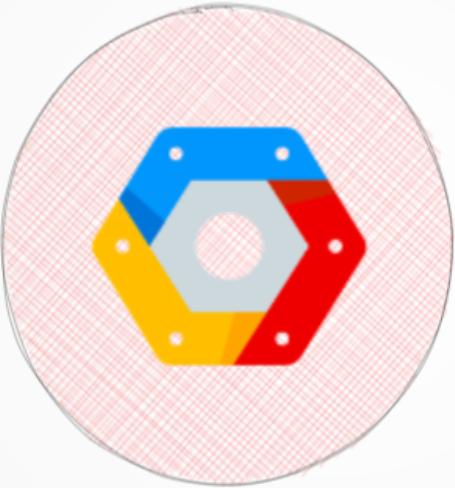
Deploy and Scale Containers on  
Managed Kubernetes



AKS

## Google Kubernetes Engine (GKE)

Managed, Production-Ready Environment  
for Containerized Applications



GKE

## Elastic Kubernetes Service (EKS)

Managed Container Service to Run  
and Scale Containerized Applications



EKS

- Serverless Kubernetes Service
- Integrated CI/CD Experience
- Enterprise Grade Security and Governance
- Co-Founder of Kubernetes Brendan Burns
- Control Plane is Free

- Fully-Managed by Google
- Running Anthos On-Premises
- Modes - Autopilot & Standard
- Control Plane - \$0.10 per Cluster/Hour

- Managed Kubernetes Service
- Fully-Managed and Integrated with AWS
- VPC Support
- Spot Instances
- Control Plane - \$0.10 per Cluster/Hour

# IaC for Data Engineers

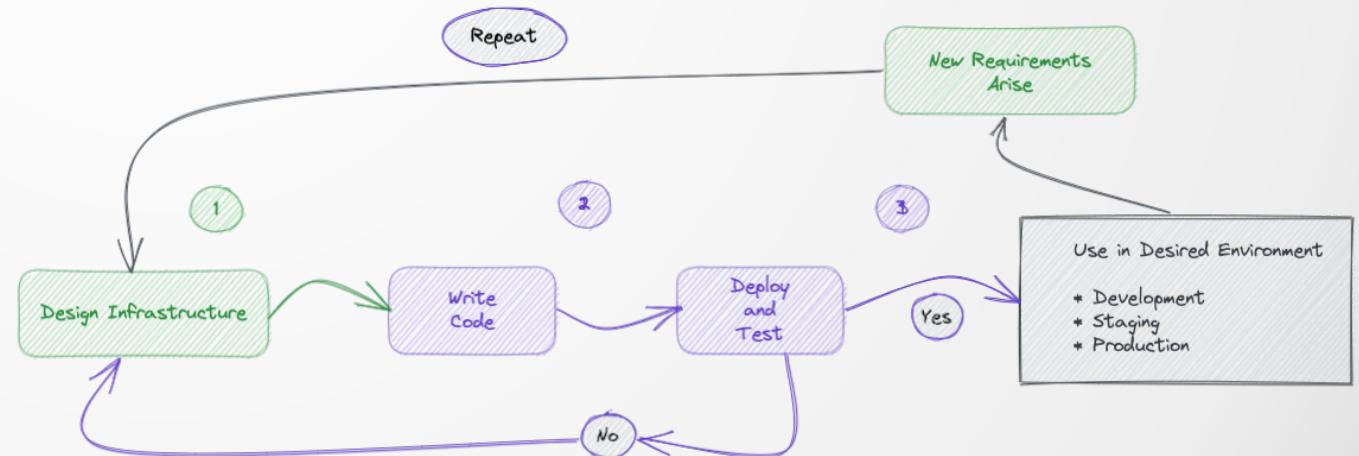
Manage and Provision Resources using Code to Reduce Manual Process, Immutable Infrastructure Approach



- \* Define Infrastructure as Code using a Declarative Language
- \* Define, Deploy, Update & Destroy your Infrastructure
- \* DevOps Key Insight = Managed Process using Code

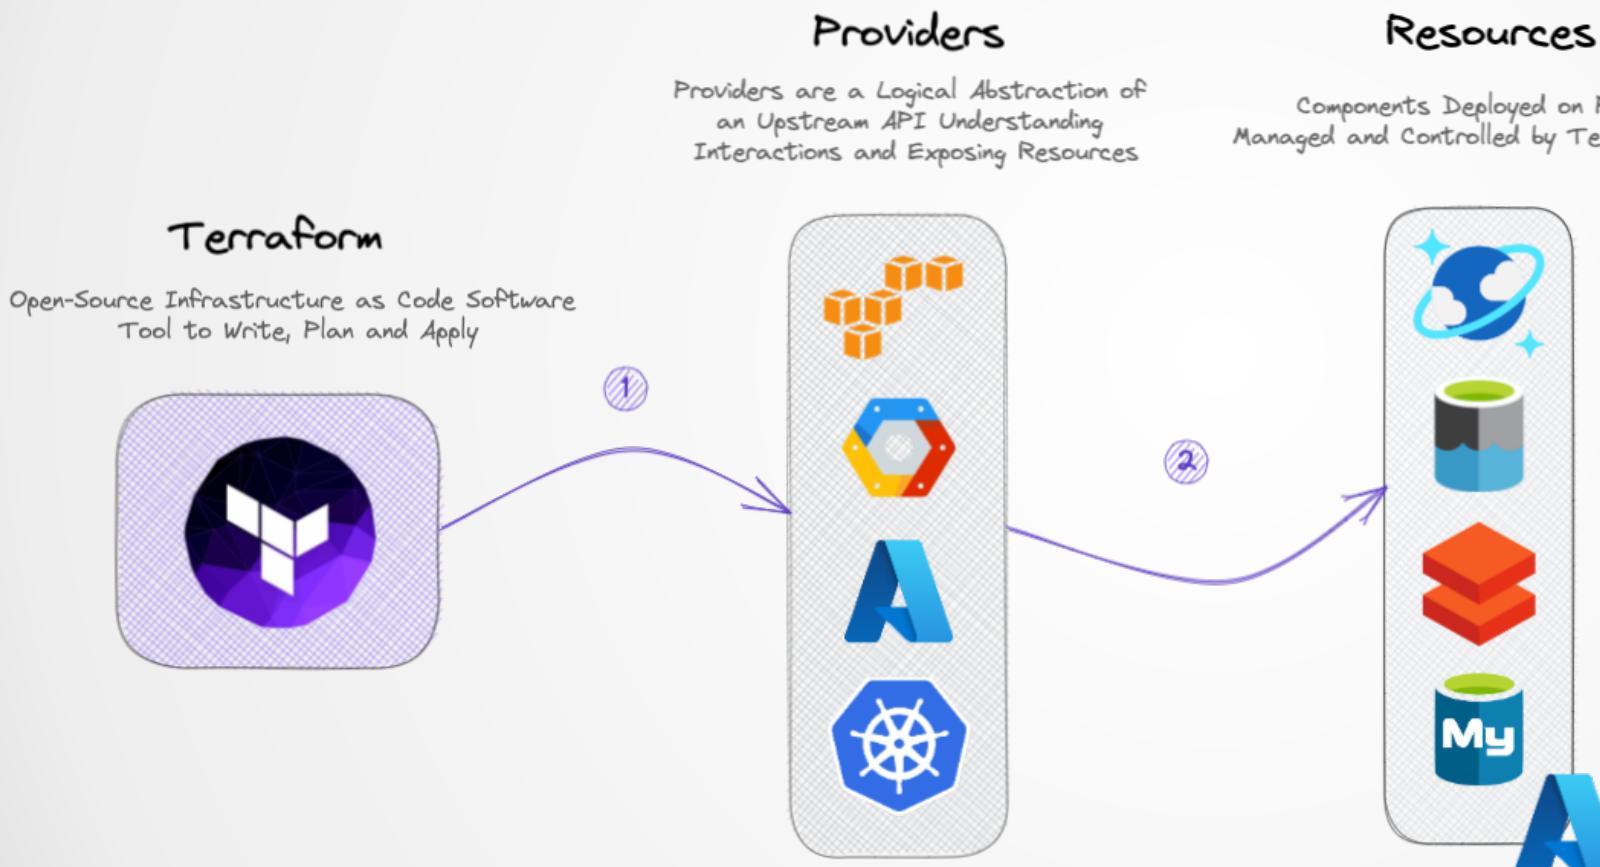
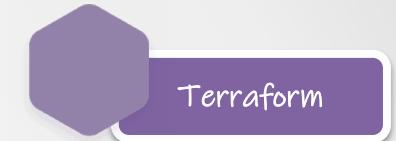


- Self-Service
- Speed and Safety
- Documentation
- Version Control
- Validation
- Reuse

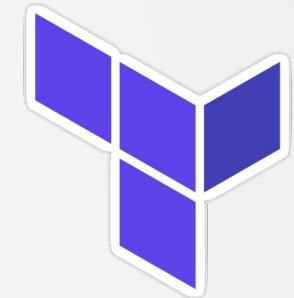


# Terraform

Open-Source Infrastructure as Code Software. Provision Resources using a Declarative Configuration Language



- Created by HashiCorp
- Initial Release 2014
- Written in Go



## Terraform Registry

Logical Abstraction of an Upstream API.  
Interacting with Providers to Expose Resources

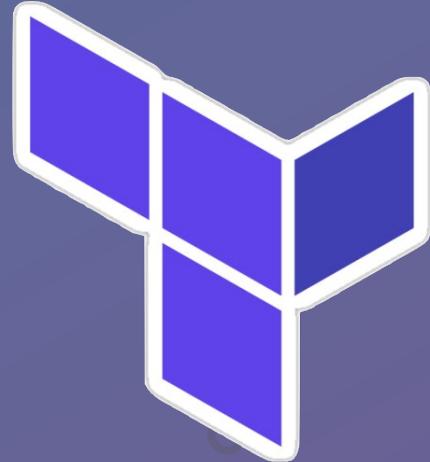
- Official
- Verified
- Community



<https://registry.terraform.io/browse/providers>

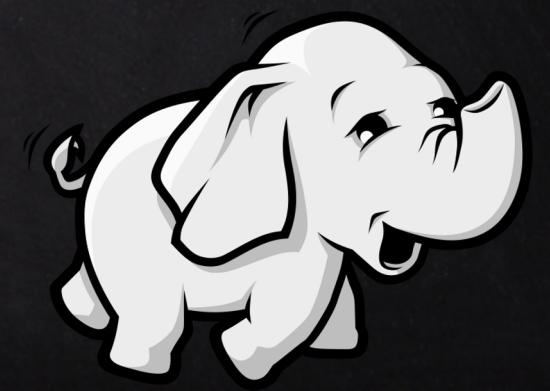
*Demo*

# Deploying Managed Kubernetes Services using Terraform



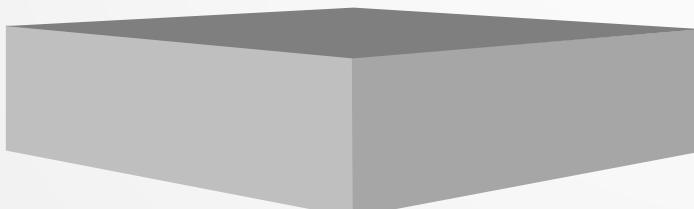
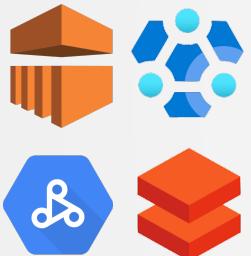
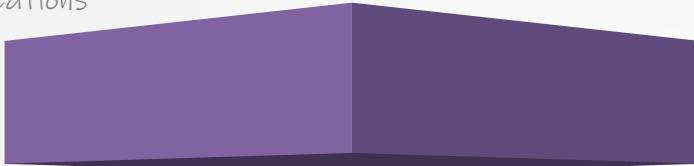
# *Debuting Kubernetes for Data Engineers*

- Big Data Eras
- Kubernetes as Cornerstone Solution for Big Data Infrastructure
- Symbiosis of Kubernetes and Big Data Systems



# Big Data Eras

On-Premises, Cloud Computing with PaaS and SaaS, Kubernetes as Cornerstone for Big Data Applications



Kubernetes as Cornerstone Solution for Apps & Data ~ [2018]



- Multi-Cloud Strategy First
- Use of Managed Kubernetes Solutions with 99.99%
- Simplify Operations By using Kubernetes as Infrastructure Backbone
- Microservice Driven Approach ~ Containerize Environments
- Cheapest Solution for Application, DataStore & Big Data Solution

Cons ~ Steep Learning Curve, Mindset Change

Use of PaaS & SaaS Software's ~ [2014]



- Solution Managed by Cloud Vendor
- Big Data as a Service ~ BDaaS
- Decoupling Computation from Storage
- Pay for Use Only
- Strategy to Reduce Overall Big Data Costs

Cons ~ Promise of Cost Reduction, Myriad of Options



Big Data on Premises using Vendors ~ [2009]

- Cloudera, MapR & HortonWorks
- Based on Open-Source Technologies
- On-Premises Data Centers using Physical Hardware
- Expensive & Dependency Wise
- Large Operations Team ~ Manage Clusters

Cons ~ Expensive, High-Level of Complexity, Burdensome

# Kubernetes as Cornerstone Solution for Big Data Infrastructure

Third Generation of Big Data Deployments, Main Characteristics and Driving Points



Cost-Wise

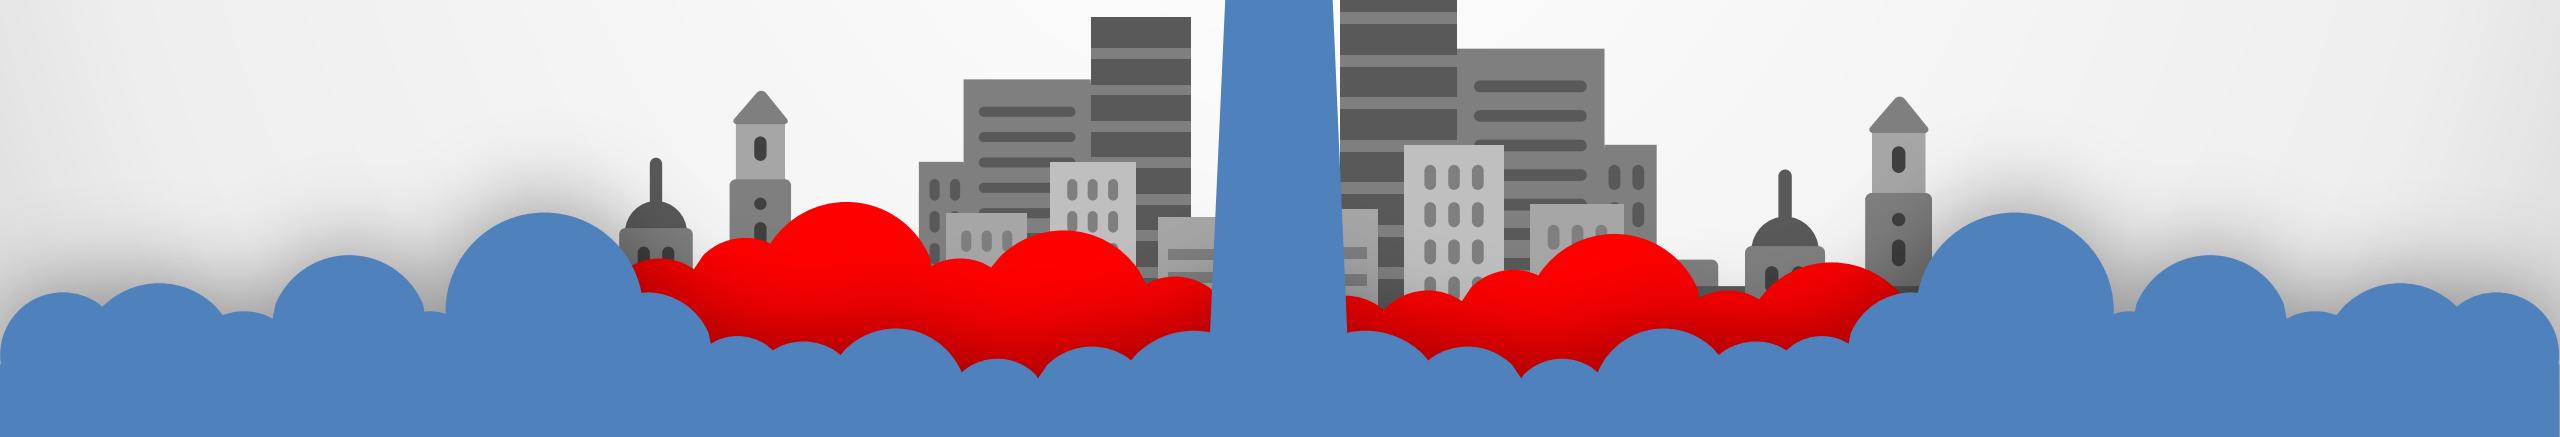


Statefulsets

Current Gen of Big Data

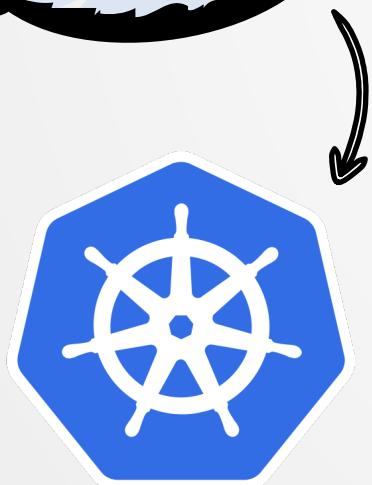
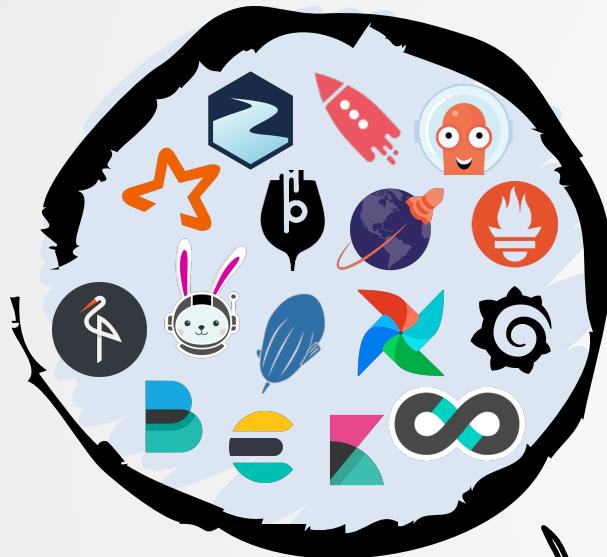


Microservices



# Symbiosis of Kubernetes and Big Data Systems

Main Components that Integrates Kubernetes and Big Data Systems  
Together in One-Single Unit



CRDs (Custom Resource) Operator Pattern



Applying State using Storage Solutions



DoK - Data on Kubernetes Community

## Custom Resources

- Software Extension
- Manage Applications & Components
- Principles of Control Loop
- Human Operator Intelligence

## Automate Repeatable Tasks

## StatefulSets API

- Used for Distributed Applications
- Ordered Pod Creation
- Stable Network Identities
- Stable Storage using Storage Class
- PV and PVCs Claims

## Sticky Identification of Pods

## Data on Kubernetes Community

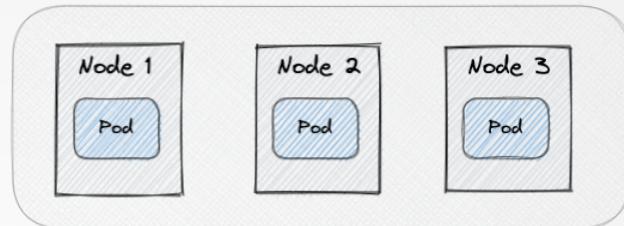
- Apply Best Practices
- Weekly Meetups
- Discussions

## Big Data Channel

<https://dokcommunity.slack.com/>

# SC, PV & PVC

Storage Class, Persistent Volume and Persistent Volume Claim for Stateful Applications



## Persistent Volume Claim (PVC)

### Persistent Volume Claim - PVC

- Pods Main Entry Communication
- Claim Space to PV

Request for Storage by User

Persistent Volume Claim (PVC)

Persistent Volume Claim (PVC)

Persistent Volume Claim (PVC)

- \* Request for Storage by User
- \* PVCs Consume PV Resources
- \* Claim Size for Volume Mount



## Persistent Volume (PV)

### Persistent Volume - PV

- Volume Plugins
- Lifecycle Independent
- Main Entry for Pod Requests

Abstracts Storage Layer

Persistent Volume (PV)

Persistent Volume (PV)

Persistent Volume (PV)

- \* Abstracts Storage Details
- \* Volume Plugins with Lifecycle Independent



## Storage Class

### Storage Class - SC

- Provide QoS Levels
- Backup Policies
- Concept of Profiles

Administrators Describe Classes

### Storage Class (SC)

#### Default

HDD

#### Premium

SSD



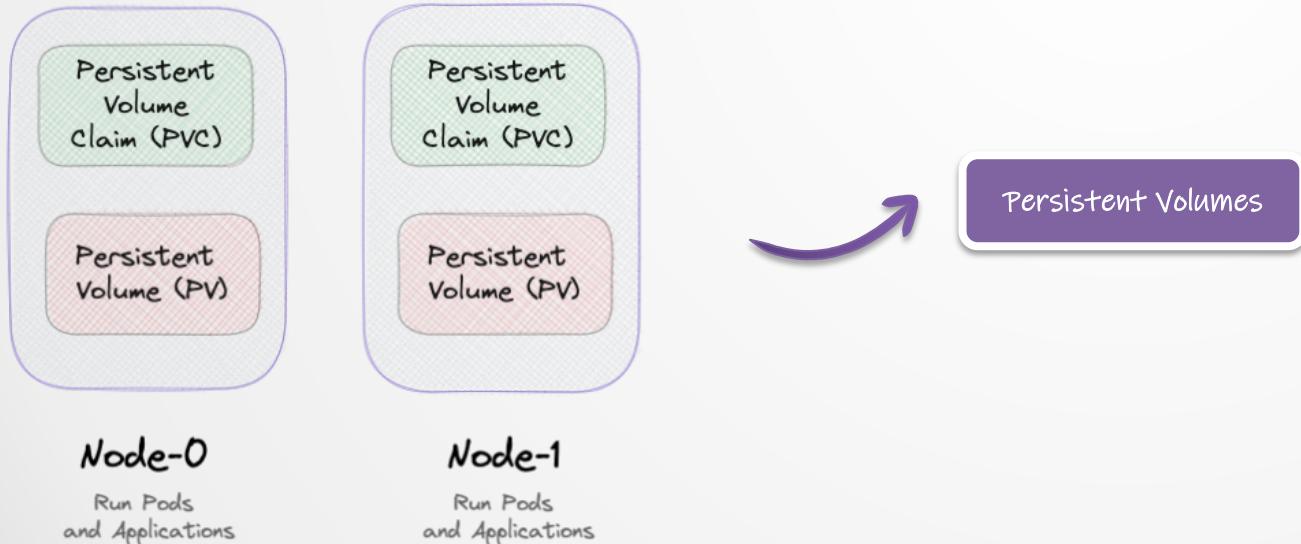
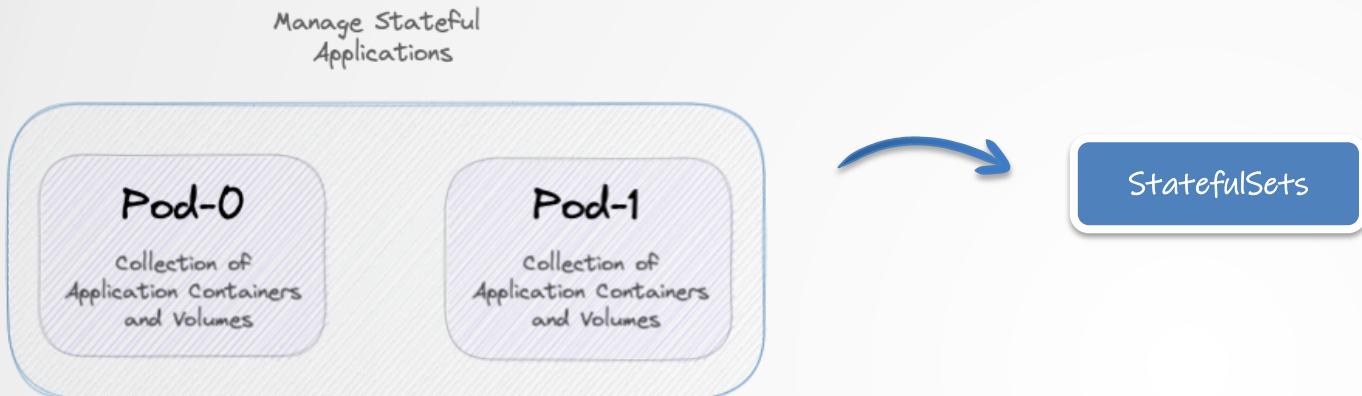
Directory Determined  
by a Volume Type

- \* AWS EBS
- \* Azure Disk
- \* Google Persistent Disk
- \* Container Storage Interface

- \* Describe Classes of Storage
- \* Affect Quality-of-Service Levels
- \* Determined by Cluster Administrators

# StatefulSets

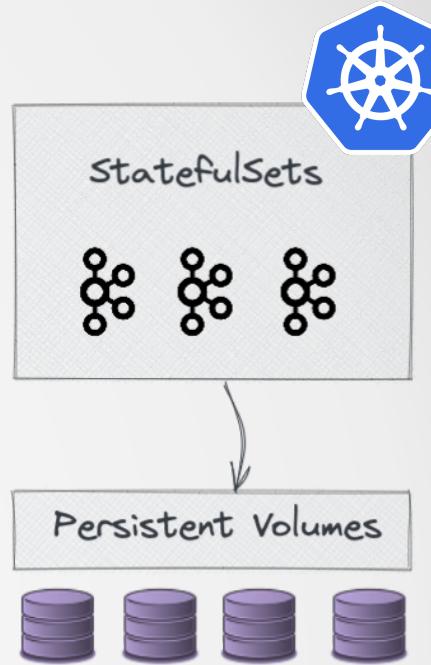
Manage Stateful Applications on Kubernetes. Provides Guarantees about Ordering and Uniqueness of Pods



## StatefulSets Explained

- Guarantees Ordering and Uniqueness of Pods
- Sticky Identity of Each Pod
- Use Deployment Type with Persistent Volumes
- Stable - Unique Network and Storage
- Ordered - Graceful Deployment and Rolling Updates

Used for Distributed Systems Deployment



## PVC & PV Explained

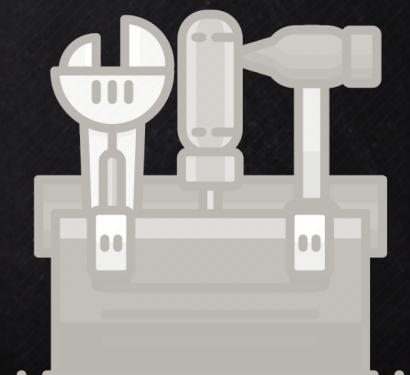
Pod Request Claim using Persistent Volume Claim API  
Persistent Volume (PV) Replies to Request from PVC  
PV Sends Requests to Storage System

Maintain Data Persisted



# *Development Toolkit for Data Engineers*

- Development Environment
- Kubernetes Package Options
- Source-Code Repository
- Container Registry Repository
- Toolbelt for Data Engineers



# Development Environment

Local and Cloud Development Options for Data Engineers

Quickly Sets Up a Local Kubernetes Cluster on  
MacOS, Linux, and Windows. Helping Developers  
for Fast Application Shipment



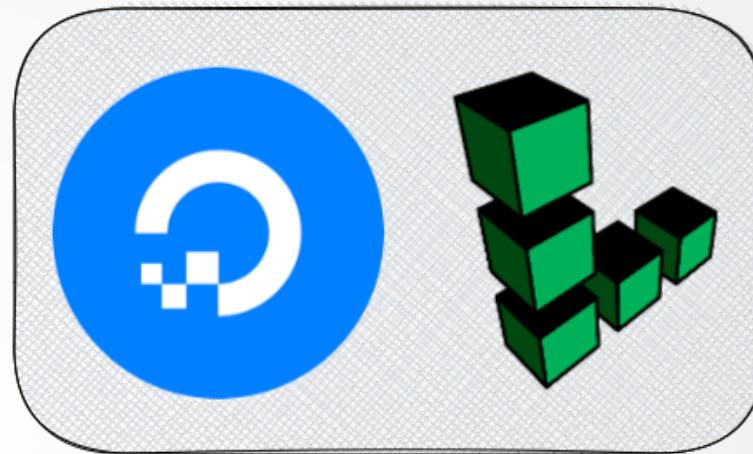
- MicroK8S
- Kind
- K3D
- 2 CPUs
- 2 GB of Memory
- 20 GB Space

Container or Virtual Machine Manager

Docker, Hyperkit, Hyper-V, KVM, Parallels,  
Podman, VirtualBox, VMware

Affordable Cloud Computing Services with  
Developer-Friendly Features and Scalability

- \* Digital Ocean
- \* Linode



Digital Ocean

- American Cloud Infrastructure
- Providers PaaS and SaaS Options
- Digital Kubernetes Engine (DKE)

<https://try.digitalocean.com/freetrialoffer/>

Linode

- American Privately-Owned Cloud Hosting
- Top Rated of 2021
- Simple, Affordable and Accessible
- Linode Kubernetes Engine (LKE)

<https://www.linode.com/p/free-credit-100/>



# Demo

## Development Environment using Minikube and Digital Ocean

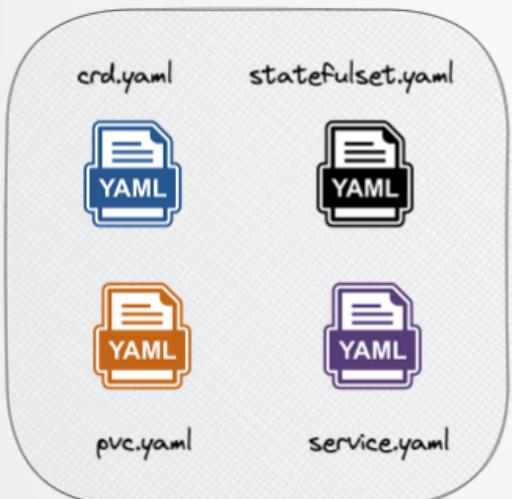


# Kubernetes Packages

Manage Complex Manifests using Tools to Reduce Complexity and Enhancing Development Experience

## Helm

Package Manager for Kubernetes. Manage Applications using Charts. Define, Install and Upgrade Complex Manifests

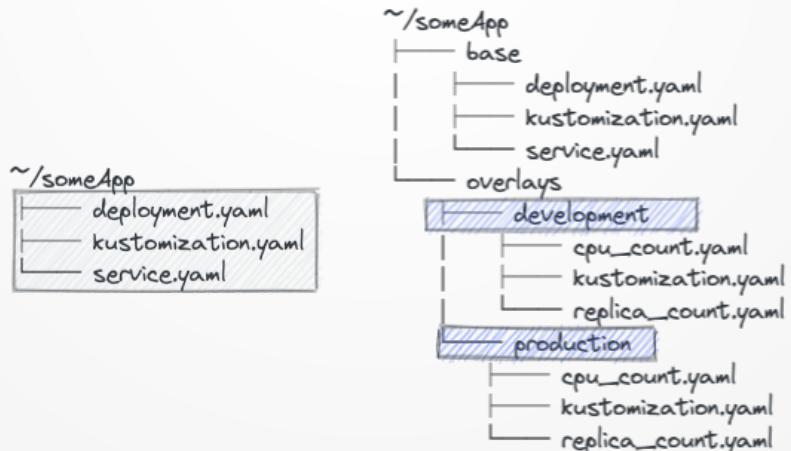


## YAML Manifests

Human-Readable text-Based Format to Easily Specify Configuration-Type Information by using Maps of Name-Value Pairs and Lists of Items

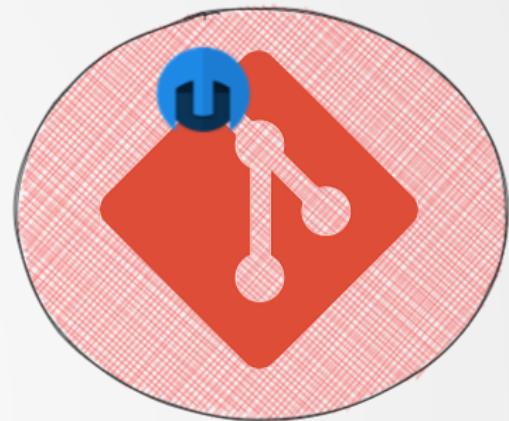
## Kustomize

Native Configuration Management using a Template-Free to Customize and Simplify Out-of-the-Shelf Applications



## ArtifactHub

Find, Install and Publish Kubernetes Packages  
Web-Based Application for CNCF Projects  
Helm Charts and Plugins



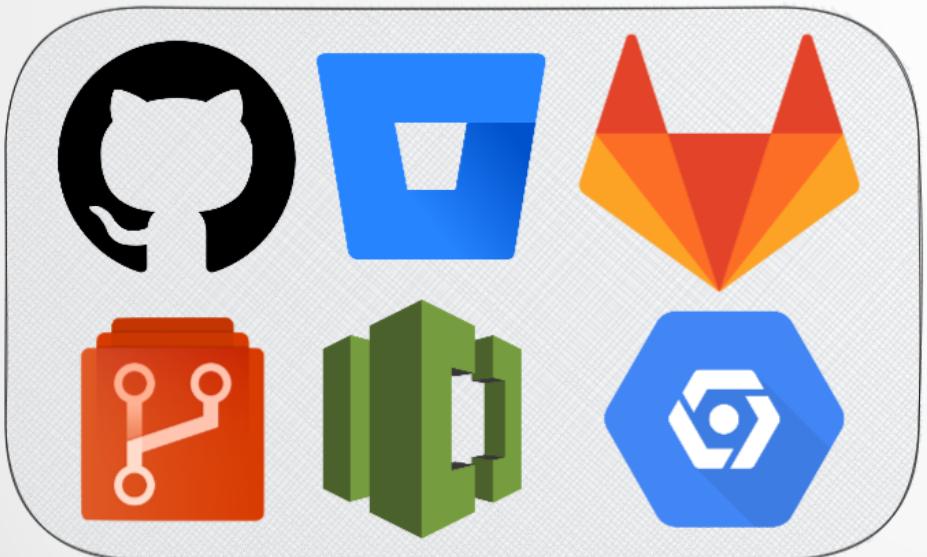
# Source-Code & Container Registry Repository

Store Deployment Files and Containerized Applications Securely in a Repository  
for Sharing and to Speed Up Development Cycle

## Source-Code Repository

Repository is a Data Structure that Stores Metadata for a Set of Files or Directory Structure

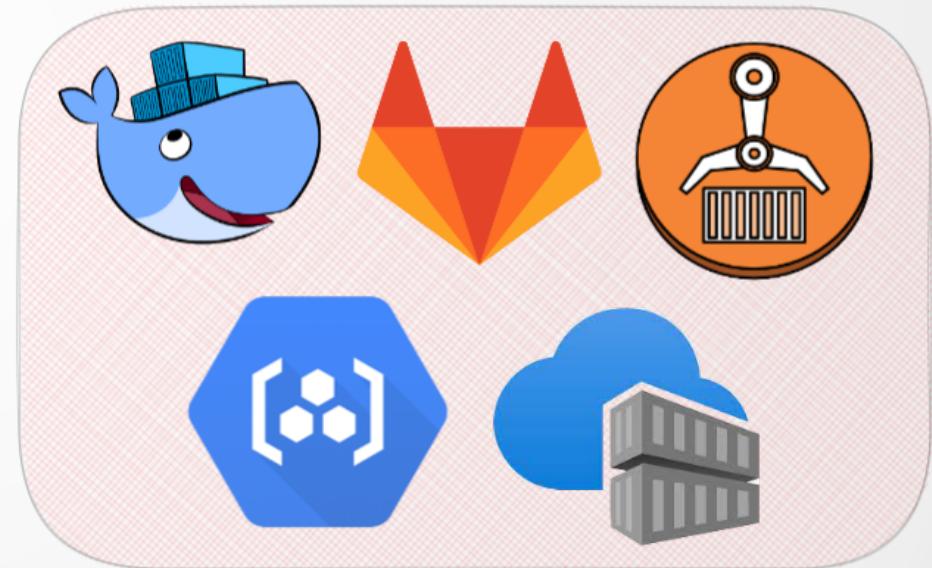
- \* GitHub
- \* BitBucket
- \* GitLab
- \* Azure Repos
- \* AWS CodeCommit
- \* Google Cloud Source



## Container Registry

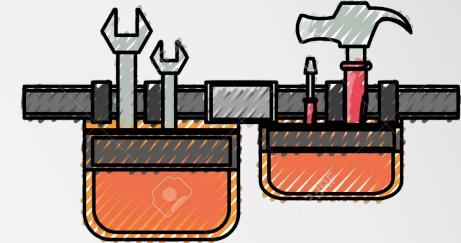
Repository or Collection Used to Store Container Images for Kubernetes, DevOps and Container-Based Application Development

- \* DockerHub
- \* GitLab Container Registry
- \* Amazon ECR
- \* Google Container Registry
- \* Azure Container Registry



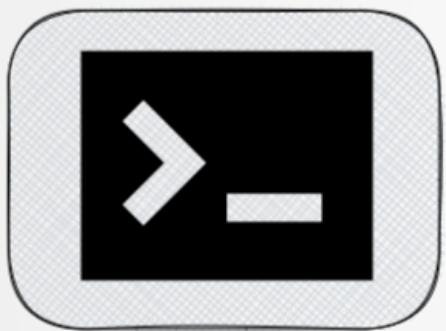
# Toolbelt for Data Engineers

Different Tools and Applications to Help Data Engineers to Maintain and Monitor Applications inside of Kubernetes



## Command-Line Interface

Operate Software and OS with Single Commands



- \* KubeCTL
- \* KubeCTX
- \* Kubens

## Log and Troubleshooting

Resolve Problems Effortlessly using Tools to Reduce Complexity



- \* Stern

## Administration & Maintenance

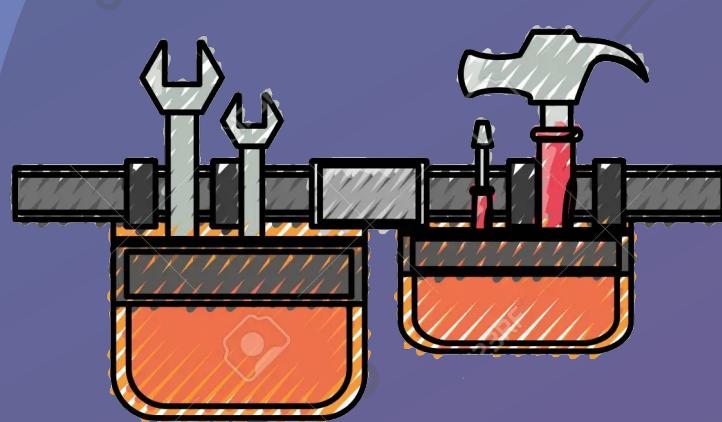
Monitor and Watch Kubernetes Cluster using best of breed tools



- \* Kube Forwarder
- \* K9S
- \* Lens

# Development Toolkit for Data Engineers

*Demo*



**If you can't fly then run, if you can't  
run then walk, if you can't walk  
then crawl, but whatever you do  
you have to keep moving forward.**

Martin Luther King Jr.